


# Introduction to Neural Networks and Deep Learning



# Agenda

## 1. Neural Network Building Blocks

- What is a Neuron?
- Working of a Neuron
- Analogy with Human Brain
- Perceptron

## 2. Why we need Neural Networks

## 3. Working of a Neural Network


- Forward Propagation
- Backward Propagation

## 4. Deep Neural Networks

- Hidden Layers
- Why Deep Neural Networks?
- Types of DNNs
  - Multilayer Perceptron (ANN)
  - Convolutional Neural Network
  - Recurrent Neural Network
  - Generative Adversarial Networks (GANs)

## 5. Applications of Deep Neural Networks

# 1. Artificial Neural Network (ANNs)



- Composed of a large number of **highly interconnected processing elements** called neurons
- Like people, **learn by example**
- Configured for **specific applications**
- Like **pattern recognition or data classification**

# Agenda

## 1. Neural Network Building Blocks

- What is a Neuron?
- Working of a Neuron
- Analogy with Human Brain
- Perceptron

## 2. Why we need Neural Networks

## 3. Working of a Neural Network


- Forward Propagation
- Backward Propagation

## 4. Deep Neural Networks

- Hidden Layers
- Why Deep Neural Networks?
- Types of DNNs
  - Multilayer Perceptron (ANN)
  - Convolutional Neural Network
  - Recurrent Neural Network
  - Generative Adversarial Networks (GANs)


## 5. Applications of Deep Neural Networks

# What is a Neuron?



- A device with many inputs and one output
- Also called McCulloch and Pitts model
- Effect that each input has at decision making depends on the weight of the particular input

# Can You Suggest the Final Outcome?



# Agenda

## 1. Neural Network Building Blocks

- What is a Neuron?
- **Working of a Neuron**
- Analogy with Human Brain
- Perceptron

## 2. Why we need Neural Networks

## 3. Working of a Neural Network


- Forward Propagation
- Backward Propagation

## 4. Deep Neural Networks

- Hidden Layers
- Why Deep Neural Networks?
- Types of DNNs
  - Multilayer Perceptron (ANN)
  - Convolutional Neural Network
  - Recurrent Neural Network
  - Generative Adversarial Networks (GANs)

## 5. Applications of Deep Neural Networks


# Working of a Neuron



- First, a neuron computes the matrix product of the weights and inputs
- It also adds a bias to the above term
- Second, it uses an activation function to get the output




# Working of a Neuron - Summation



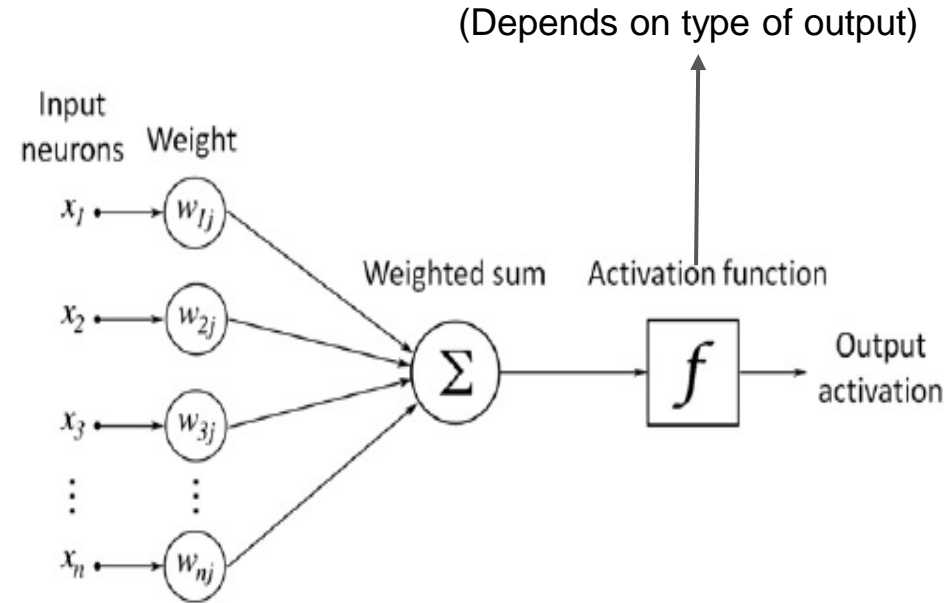
- Firstly, the inputs are matrix multiplied to the weight vector
$$\begin{bmatrix} Input1 & Input2 \end{bmatrix} * \begin{bmatrix} Weight1 \\ Weight2 \end{bmatrix}$$
- Size of the weight vector = Number of inputs
- Then, bias is added to get the output 'z'

# Working of a Neuron - Matrix Operations




- Ensure that **dimensions of input and weights** are set up for matrix multiplication
- **Bias is broadcasted** to be added to each row

# Working of a Neuron - Activation




- An **activation is applied to output 'z'** to compute final output 'a' of a neuron
- Different **activations modify the output in different ways**

# Working of a Neuron - Sigmoid Activation Function




- Sigmoid activation is a nonlinear function
- It exponentially transforms the values the further away they are from 0
- We expect to get a non-zero gradient between the red lines

# Working of a Neuron - ReLU Activation Function



- Rectified Linear Unit or ReLU is also a nonlinear function
- It nullifies inputs less than zero to zero
- The inputs greater than zero are kept same

# Working of a Neuron - tanh Activation Function



- Hyperbolic tangent or  $\tanh$  is another nonlinear activation
- Similar to sigmoid except that for negative inputs, it gives negative output

\*Details of the activation functions will be covered in later sessions

# Agenda

## 1. Neural Network Building Blocks

- What is a Neuron?
- Working of a Neuron
- **Analogy with Human Brain**
- Perceptron

## 2. Why we need Neural Networks

## 3. Working of a Neural Network

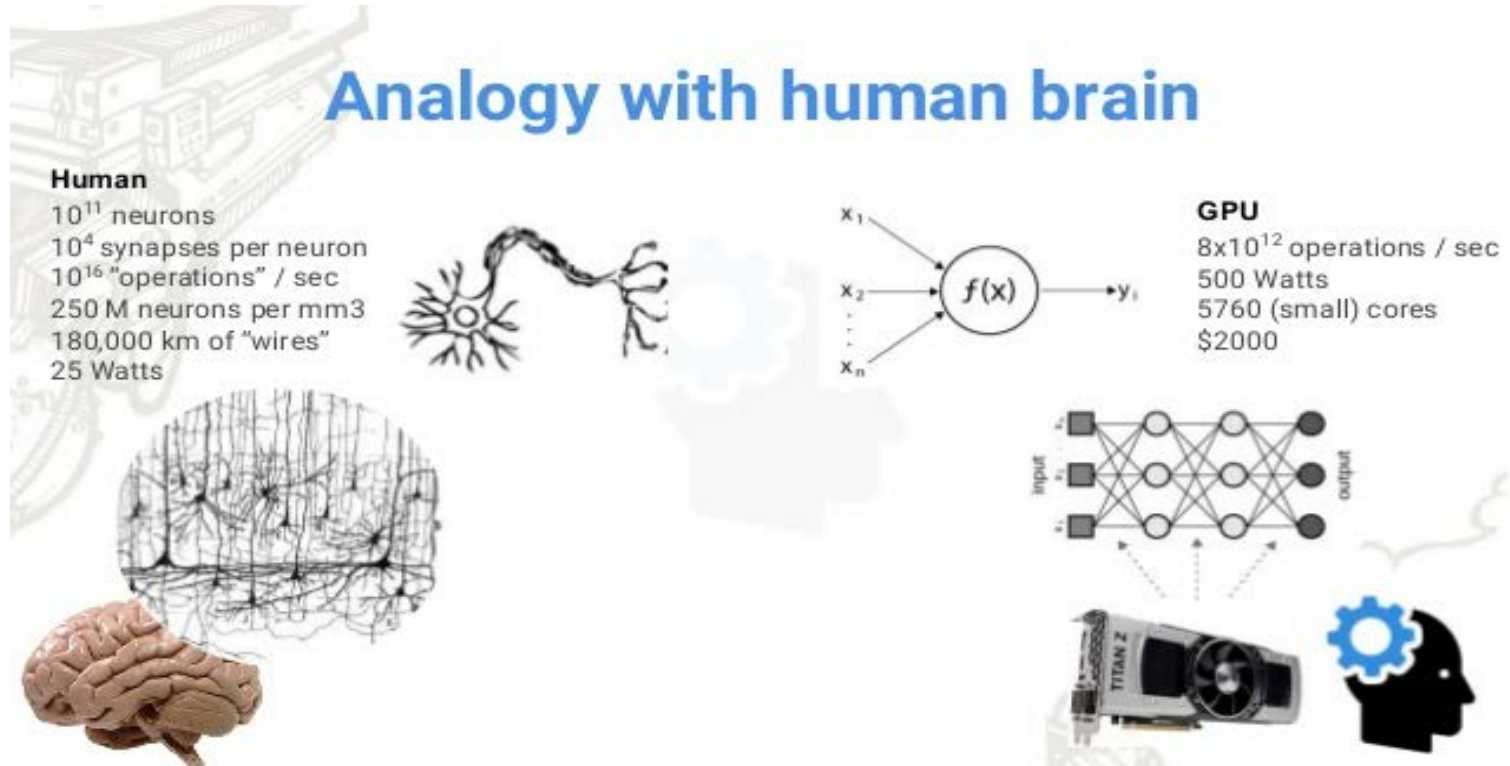
- Forward Propagation
- Backward Propagation

## 4. Deep Neural Networks

- Hidden Layers
- Why Deep Neural Networks?
- Types of DNNs
  - Multilayer Perceptron (ANN)
  - Convolutional Neural Network
  - Recurrent Neural Network
  - Generative Adversarial Networks (GANs)


## 5. Applications of Deep Neural Networks

# Analogy with Human Brain






# How Human Brain Learns




- A typical neuron **collects signals through** a host of fine structures called **dendrites**
- Neuron sends out **spikes of electrical activity** through a **long thin strand known as an axon**
- Axons split into **thousands of branches**

# How Human Brain Learns




- At the end of each branch, a *synapse* converts activity from axon into electrical effects
- These inhibit or excite activity in the connected neurons
- Learning occurs by changing the effectiveness of the synapses so that the influence of one neuron on another changes.


# Analogy of Brain to Neural Network



This image by Felipe Peruchio is licensed under [CC-BY 3.0](#)



# Neural Network terminology in a Neuron



- We first try to deduce **essential features of neurons** and their interconnections
- Since our **knowledge of neurons is incomplete and computing power is limited**
- Our models are actually gross **idealisations of real networks of neurons**

# Agenda

## 1. Neural Network Building Blocks

- What is a Neuron?
- Working of a Neuron
- Analogy with Human Brain
- **Perceptron**

## 2. Why we need Neural Networks

## 3. Working of a Neural Network

- Forward Propagation
- Backward Propagation


## 4. Deep Neural Networks

- Hidden Layers
- Why Deep Neural Networks?
- Types of DNNs
  - Multilayer Perceptron (ANN)
  - Convolutional Neural Network
  - Recurrent Neural Network
  - Generative Adversarial Networks (GANs)

## 5. Applications of Deep Neural Networks

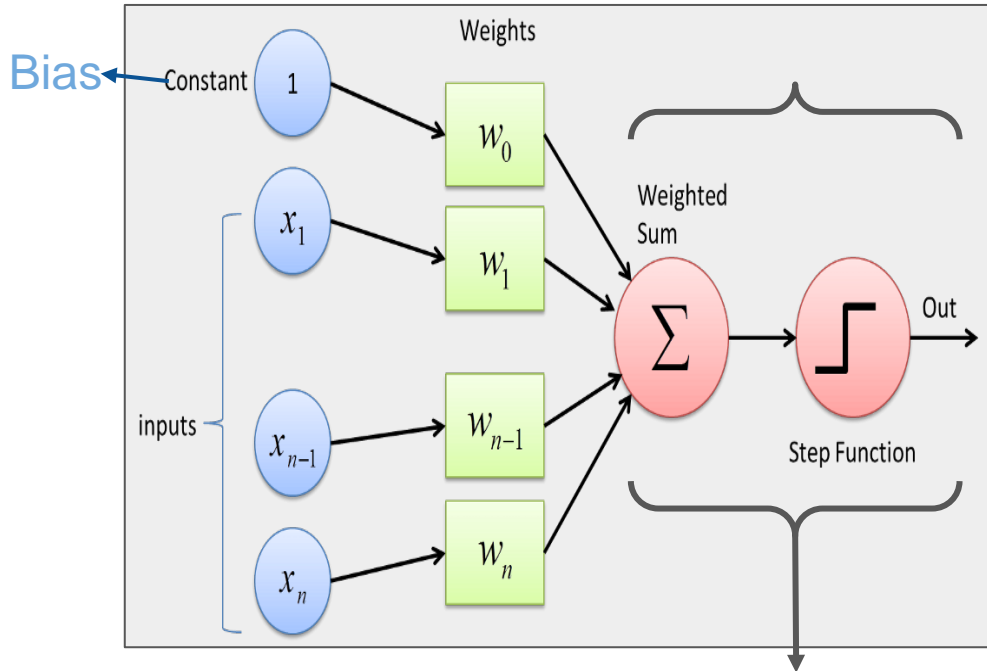
# Perceptron - Introduction

## SLP Architecture



- A perceptron is commonly taken as a **Single Layer Perceptron (SLP)**
- This means that there is only a **single layer of neurons**
- The **number of neurons in the SLP** may vary according to the input


# Perceptron



Single neuron of the layer.  
Multiple such neurons can exist in an  
SLP

- Perceptron is a **more general computational model** than McCulloch-Pitts neuron
- It takes an **input and aggregates it** (weighted sum)
- Returns **1 only if the aggregated sum is more than some threshold**, else returns 0

# Perceptron - Working



A more accepted convention,

$$y = 1 \quad \text{if} \quad \sum_{i=0}^n w_i * x_i \geq 0$$

$$= 0 \quad \text{if} \quad \sum_{i=0}^n w_i * x_i < 0$$

where,  $x_0 = 1$  and  $w_0 = -\theta$



# Agenda

## 1. Neural Network Building Blocks

- What is a Neuron?
- Working of a Neuron
- Analogy with Human Brain
- Perceptron

## 2. Why we need Neural Networks

## 3. Working of a Neural Network


- Forward Propagation
- Backward Propagation

## 4. Deep Neural Networks


- Hidden Layers
- Why Deep Neural Networks?
- Types of DNNs
  - Multilayer Perceptron (ANN)
  - Convolutional Neural Network
  - Recurrent Neural Network
  - Generative Adversarial Networks (GANs)

## 5. Applications of Deep Neural Networks

# Why do we need Neural Networks - Example




# Classification Example - Linear Classifier



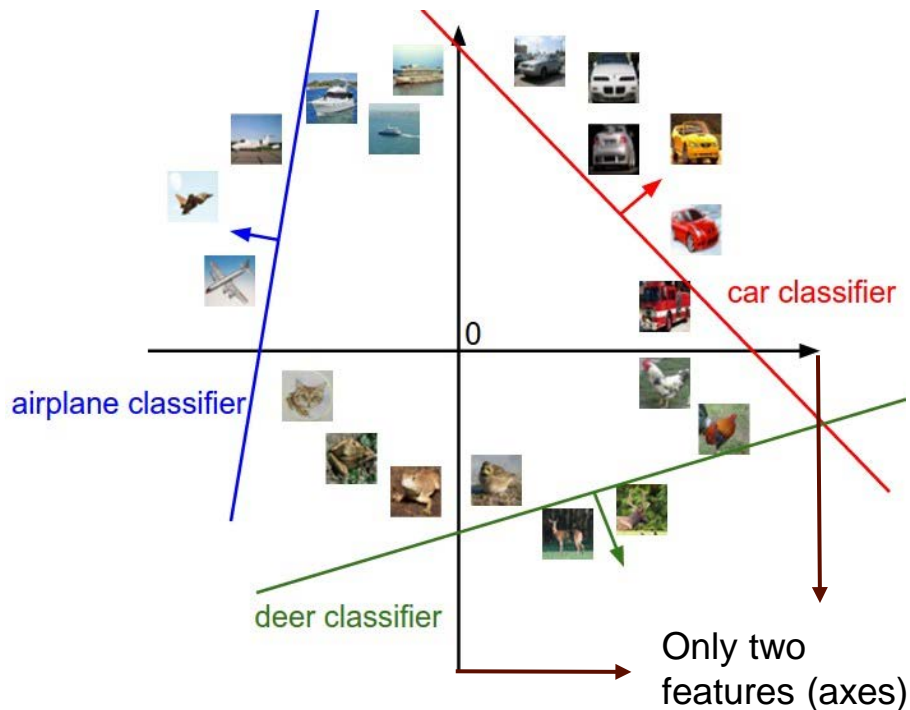
- Linear classifier - fails miserably
- Boundary of separation - meant to be nonlinear
- Regular machine learning algorithms won't work

# Classification Problem - Neural Network Classifier



- Neural networks have the ability to build *arbitrary shaped classification boundaries*
- This difference becomes even more evident when there are more than two classes

# Multiclass Classification Problem



- Linear classifiers face even **more problems for multiclass classification**
- **256x256 sized image -  $2^{16}$  features**
- Neural Networks **classify multi-outputs on large no. of features and classes**

# Agenda

## 1. Neural Network Building Blocks

- What is a Neuron?
- Working of a Neuron
- Analogy with Human Brain
- Perceptron

## 2. Why we need Neural Networks

## 3. Working of a Neural Network

- Forward Propagation
- Backward Propagation


## 4. Deep Neural Networks

- Hidden Layers
- Why Deep Neural Networks?
- Types of DNNs
  - Multilayer Perceptron (ANN)
  - Convolutional Neural Network
  - Recurrent Neural Network
  - Generative Adversarial Networks (GANs)

## 5. Applications of Deep Neural Networks



# Neural Networks - Working



Output of one layer used as input to next  
Hidden layers > 1 => Multi-layer Perceptron

- One perceptron = one decision
- What about multiple decisions? ◦ E.g. Multi-image classification
- Stack as many output nodes as the number of possible outcomes into the final layer
  - Neural network

# Agenda

## 1. Neural Network Building Blocks

- What is a Neuron?
- Working of a Neuron
- Analogy with Human Brain
- Perceptron

## 2. Why we need Neural Networks

## 3. Working of a Neural Network

- Forward Propagation
- Backward Propagation


## 4. Deep Neural Networks

- Hidden Layers
- Why Deep Neural Networks?
- Types of DNNs
  - Multilayer Perceptron (ANN)
  - Convolutional Neural Network
  - Recurrent Neural Network
  - Generative Adversarial Networks (GANs)

## 5. Applications of Deep Neural Networks




# Forward Propagation




- Let's take student study-sleep dataset as an example
- At each layer, compute matrix product and its activation, pass to the next layer as input

# Forward Propagation



- Two features => two input nodes
- Regular neural network => single hidden layer
- Nodes in hidden layers - our choice
- Output - test score
- Need only one node

# Forward Propagation




## Variables

Math Symbol	Definition	Dimensions	Example
$X$	Input Data, each row in an example	$(\text{numExamples}, \text{inputLayerSize})$	$(N, 2)$
$y$	target data	$(\text{numExamples}, \text{outputLayerSize})$	$(N, 1)$
$W^{(1)}$	Layer 1 weights	$(\text{inputLayerSize}, \text{hiddenLayerSize})$	$(2, 3)$
$W^{(2)}$	Layer 2 weights	$(\text{hiddenLayerSize}, \text{outputLayerSize})$	$(3, 1)$
$z^{(2)}$	Layer 2 activation	$(\text{numExamples}, \text{hiddenLayerSize})$	$(N, 3)$
$a^{(2)}$	Layer 2 activity	$(\text{numExamples}, \text{hiddenLayerSize})$	$(N, 3)$
$z^{(3)}$	Layer 3 activation	$(\text{numExamples}, \text{outputLayerSize})$	$(N, 1)$


N - No. of training examples

# Forward Propagation



- Example matrix of 3 instances, each of which has an hours slept and hours studied value
- Eg. instance 1 - 3 hours of sleep and 5 hours of study
- Predict the test scores for the three instances

# Forward Propagation



- These are the **main equations** which will be used to calculate the output
- We **ignore the bias** here for simplicity
- We do the **same steps progressively for each layer**

Matrix product at layer 2 -  $z^{(2)} = XW^{(1)}$

Activation at layer 2 -  $a^{(2)} = f(z^{(2)})$

Matrix product at layer 3 -  $z^{(3)} = a^{(2)}W^{(2)}$

Activation at layer 3 (output) -  $\hat{y} = f(z^{(3)})$


# Forward Propagation - Matrix Multiplication (1st layer)

For first layer -  $z^{(2)} = XW^{(1)}$

$$= \begin{bmatrix} 3 & 5 \\ 5 & 1 \\ 10 & 2 \end{bmatrix}_{3 \times 2} \begin{bmatrix} W_{11}^{(1)} & W_{12}^{(1)} & W_{13}^{(1)} \\ W_{21}^{(1)} & W_{22}^{(1)} & W_{23}^{(1)} \end{bmatrix}_{2 \times 3}$$

- Multiply 3 x 2 input matrix with 2 x 3 weight matrix
- For eg.  $W_{23}^{(1)}$  means the weight of the edge from the second (2) input feature to the third (3) node of the first (1) hidden layer

# Forward Propagation




$$z^{(2)} = XW^{(1)} = \begin{bmatrix} 3 & 5 \\ 5 & 1 \\ 10 & 2 \end{bmatrix} \begin{bmatrix} W_{11}^{(1)} & W_{12}^{(1)} & W_{13}^{(1)} \\ W_{21}^{(1)} & W_{22}^{(1)} & W_{23}^{(1)} \end{bmatrix}$$

$$= \begin{bmatrix} 3W_{11}^{(1)} + 5W_{21}^{(1)} & 3W_{12}^{(1)} + 5W_{22}^{(1)} & 3W_{13}^{(1)} + 5W_{23}^{(1)} \\ 5W_{11}^{(1)} + W_{21}^{(1)} & 5W_{12}^{(1)} + W_{22}^{(1)} & 5W_{13}^{(1)} + W_{23}^{(1)} \\ 10W_{11}^{(1)} + 2W_{21}^{(1)} & 10W_{12}^{(1)} + 2W_{22}^{(1)} & 10W_{13}^{(1)} + 2W_{23}^{(1)} \end{bmatrix}$$

- Each entry in  $z$  - sum of weighted inputs to each hidden neuron.
- $z$  is  $3 \times 3$  matrix, one row for each sample, and one column for each hidden unit ( $N=3$ ,  $\rightarrow 3 \times 3$  matrix)

# Forward Propagation - Activation

- Matrix activated using an activation function
- Sigmoid activation function as example
- Calculate the sigmoid of each value in the matrix



$$a^{(2)} = \text{sigmoid} \left( \begin{bmatrix} 3W_{11}^{(1)} + 5W_{21}^{(1)} & 3W_{12}^{(1)} + 5W_{22}^{(1)} & 3W_{13}^{(1)} + 5W_{23}^{(1)} \\ 5W_{11}^{(1)} + W_{21}^{(1)} & 5W_{12}^{(1)} + W_{22}^{(1)} & 5W_{13}^{(1)} + W_{23}^{(1)} \\ 10W_{11}^{(1)} + 2W_{21}^{(1)} & 10W_{12}^{(1)} + 2W_{22}^{(1)} & 10W_{13}^{(1)} + 2W_{23}^{(1)} \end{bmatrix} \right)$$

$$= \begin{bmatrix} \text{smd}(3W_{11}^{(1)} + 5W_{21}^{(1)}) & \text{smd}(3W_{12}^{(1)} + 5W_{22}^{(1)}) & \text{smd}(3W_{13}^{(1)} + 5W_{23}^{(1)}) \\ \text{smd}(5W_{11}^{(1)} + W_{21}^{(1)}) & \text{smd}(5W_{12}^{(1)} + W_{22}^{(1)}) & \text{smd}(5W_{13}^{(1)} + W_{23}^{(1)}) \\ \text{smd}(10W_{11}^{(1)} + 2W_{21}^{(1)}) & \text{smd}(10W_{12}^{(1)} + 2W_{22}^{(1)}) & \text{smd}(10W_{13}^{(1)} + 2W_{23}^{(1)}) \end{bmatrix}$$

\* Sigmoid abbreviated as smd



# Forward Propagation - Initialization

Random Initialization

$$\begin{bmatrix} W_{11}^{(1)} & W_{12}^{(1)} & W_{13}^{(1)} \\ W_{21}^{(1)} & W_{22}^{(1)} & W_{23}^{(1)} \end{bmatrix} \longrightarrow \begin{bmatrix} 0.8 & 0.4 & 0.3 \\ 0.2 & 0.9 & 0.5 \end{bmatrix} \quad \checkmark$$


Normalized (0,1) values for easy calculation

Zero Initialization

$$\begin{bmatrix} W_{11}^{(1)} & W_{12}^{(1)} & W_{13}^{(1)} \\ W_{21}^{(1)} & W_{22}^{(1)} & W_{23}^{(1)} \end{bmatrix} \longrightarrow \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix} \quad \times$$

All computed gradients evaluate to 0


# Forward Propagation



$$\begin{aligned} a^{(2)} &= f(z^{(2)}) \\ z^{(3)} &= a^{(2)} W^{(2)} \\ \hat{y} &= f(z^{(3)}) \end{aligned}$$
$$W^{(2)} = \begin{bmatrix} W_{11}^{(2)} \\ W_{21}^{(2)} \\ W_{31}^{(2)} \end{bmatrix}$$

- $a^{(2)}$  passed as input to the third (final) layer
- Matrix multiplication -  $a^{(2)}$  (3 x 3) and  $W^{(2)}$  (3 x 1)
- Final output - activation matrix of the final layer

# Forward Propagation



- Final output - **sigmoid of the matrix  $z^{(3)}$**
- No training => **Poor output predictions** of test scores
- This is **one iteration of forward propagation**
- To get good results, we train our network i.e. **use a feedback mechanism to realign weights**

$$z^{(3)} = a^{(2)} W^{(2)}$$

$$= \begin{bmatrix} \text{smd}(3W_{11}^{(1)} + 5W_{21}^{(1)}) & \text{smd}(3W_{12}^{(1)} + 5W_{22}^{(1)}) & \text{smd}(3W_{13}^{(1)} + 5W_{23}^{(1)}) \\ \text{smd}(5W_{11}^{(1)} + W_{21}^{(1)}) & \text{smd}(5W_{12}^{(1)} + W_{22}^{(1)}) & \text{smd}(5W_{13}^{(1)} + W_{23}^{(1)}) \\ \text{smd}(10W_{11}^{(1)} + 2W_{21}^{(1)}) & \text{smd}(10W_{12}^{(1)} + 2W_{22}^{(1)}) & \text{smd}(10W_{13}^{(1)} + 2W_{23}^{(1)}) \end{bmatrix} \begin{bmatrix} W_{11}^{(2)} \\ W_{21}^{(2)} \\ W_{31}^{(2)} \end{bmatrix} \xrightarrow{f(\cdot)} \hat{y} = f(z^{(3)})$$

# Agenda

## 1. Neural Network Building Blocks

- What is a Neuron?
- Working of a Neuron
- Analogy with Human Brain
- Perceptron

## 2. Why we need Neural Networks

## 3. Working of a Neural Network


- Forward Propagation
- **Backward Propagation**

## 4. Deep Neural Networks

- Hidden Layers
- Why Deep Neural Networks?
- Types of DNNs
  - Multilayer Perceptron (ANN)
  - Convolutional Neural Network
  - Recurrent Neural Network
  - Generative Adversarial Networks (GANs)

## 5. Applications of Deep Neural Networks

# Backward Propagation (Backprop)



- Need to calculate the cost function gradient
- Requires known, target data for each input
- Supervised training


# Backward Propagation - Cost

Minimize cost function (sum of squared loss) -

$$J = \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

- Here,  $N$  is the no. of samples,  $y_i$  is the actual output,  $\hat{y}_i$  is our prediction
- Complete equations for gradient descent - too long and complex
- Since we have already dealt with the concept, we use it with basic hints

# Backward Propagation - Weights



- We have a collection of 9 weights
- We're going to make our cost ( $J$ ) as small as possible
- This can be done using an optimal combination of the weights



# Backward Propagation - Brute Force vs GD


## Brute Force

## Gradient Descent

<ul style="list-style-type: none"><li>• Check cost by <b>only tweaking the first weight</b> <math>W_{11}</math> for <b>1000 different values</b> = 0.11 secs</li></ul>	<ul style="list-style-type: none"><li>• <b>Calculate the cost function</b> using all samples</li></ul>
<ul style="list-style-type: none"><li>• Check 1000 values of <math>W_{11}</math> and <math>W_{12}</math> = <b>~100 secs</b></li></ul>	<ul style="list-style-type: none"><li>• <b>Calculate gradient</b> to update the weights = <b>~10 seconds</b></li></ul>
<ul style="list-style-type: none"><li>• Check 1000 values <b>for all 9 weights</b> = <b>1 trillion millenium</b></li></ul>	<ul style="list-style-type: none"><li>• <b>Update the weights</b> using the gradient = <b>~30 seconds</b></li></ul>
<ul style="list-style-type: none"><li>• That's a billion times more than the <b>age of the universe</b></li></ul>	



# Backward Propagation - Training (Gradient Descent)



- Smarter way - Gradient Descent
- Capable of incredible speedups in higher dimensions
- Gradient = derivative of the cost wrt the weights
- Three choices of **gradient descent**: batch(standard), stochastic or mini-batch

# Backward Propagation - Batch, Stochastic and Mini-batch GD


Batch Gradient Descent -  $\sum_{i=1}^N \frac{\partial J}{\partial W}$

Stochastic Gradient Descent -  $\frac{\partial J}{\partial W}$

Mini-batch Gradient Descent -  $\sum_{i=1}^s \frac{\partial J}{\partial W}$


- **Batch** gradient descent method **sums up all the derivatives of J for all samples (N)**
- **Stochastic** gradient descent (SGD) method **uses one derivative at one sample and move to another sample point**
- **Mini-batch** gradient descent **uses fixed batch size of samples** (here, s), usually much smaller than N

# Batch vs Stochastic vs Mini-Batch Gradient Descent



- **Batch** : takes fewer steps, high no. of calculations per step
- **Stochastic** : takes high no. of steps, only one calculation per step
- **Mini-batch** : Middle ground, takes intermediate no. of steps, few calculations per step
- Performance :  
Mini-batch > Stochastic > Batch

# Backward Propagation - Output



- Backprop uses **gradient descent to train our neural networks efficiently**
- Gives an **accurate measure** very close to the actual output

# Agenda

## 1. Neural Network Building Blocks

- What is a Neuron?
- Working of a Neuron
- Analogy with Human Brain
- Perceptron

## 2. Why we need Neural Networks

## 3. Working of a Neural Network


- Forward Propagation
- Backward Propagation

## 4. Deep Neural Networks

- Hidden Layers
- Why Deep Neural Networks?
- Types of DNNs
  - Multilayer Perceptron (ANN)
  - Convolutional Neural Network
  - Recurrent Neural Network
  - Generative Adversarial Networks (GANs)

## 5. Applications of Deep Neural Networks

# Deep Neural Networks



# Agenda

## 1. Neural Network Building Blocks

- What is a Neuron?
- Working of a Neuron
- Analogy with Human Brain
- Perceptron

## 2. Why we need Neural Networks

## 3. Working of a Neural Network

- Forward Propagation
- Backward Propagation


## 4. Deep Neural Networks

- **Hidden Layers**
- Why Deep Neural Networks?
- Types of DNNs
  - Multilayer Perceptron (ANN)
  - Convolutional Neural Network
  - Recurrent Neural Network
  - Generative Adversarial Networks (GANs)

## 5. Applications of Deep Neural Networks


# Deep Neural Networks (DNN) - Hidden Layers

Simple Neural Network



● Input Layer

Deep Learning Neural Network




● Hidden Layer    ● Output Layer

- Conventional NN - one hidden layer
- DNN - two or more hidden layers
- *Deep layers => Deep Learning*
- Previous algorithms like forward propagation, gradient descent, backward propagation etc. remain same






# Deep Neural Network - Hidden Layers



Gender Classifier example

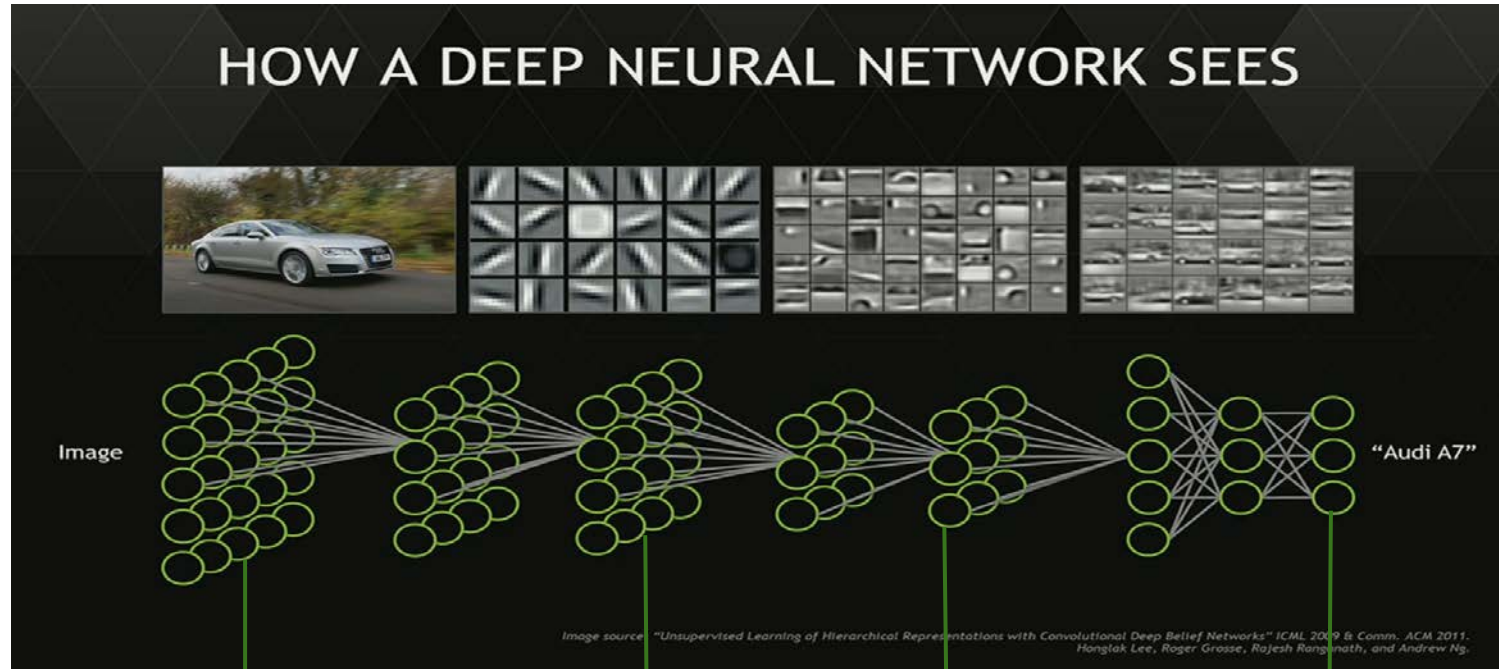


Input nodes = no.  
of features

Arbitrary no. of hidden layers and  
their no. of nodes

Output nodes =  
no. of classes

# Deep Neural Network - Hidden Layers



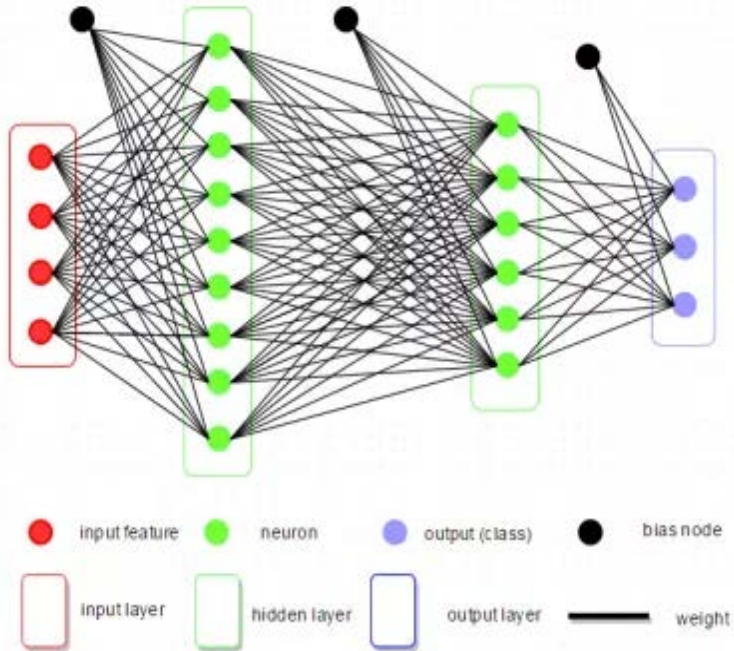
Input layer -> nodes = no.  
of features

Hidden layers capture features  
like edges, shapes, wheels, etc.

Output layer  
predicts object

# Choosing No. of layers and nodes

A 3-layers fully connected neural network (DNN)



- Generic way - Experimentation
- Develop intuition with experience
- Depth - the more the better, but with diminishing returns
- Take ideas from papers, models

# Agenda

## 1. Neural Network Building Blocks

- What is a Neuron?
- Working of a Neuron
- Analogy with Human Brain
- Perceptron

## 2. Why we need Neural Networks

## 3. Working of a Neural Network

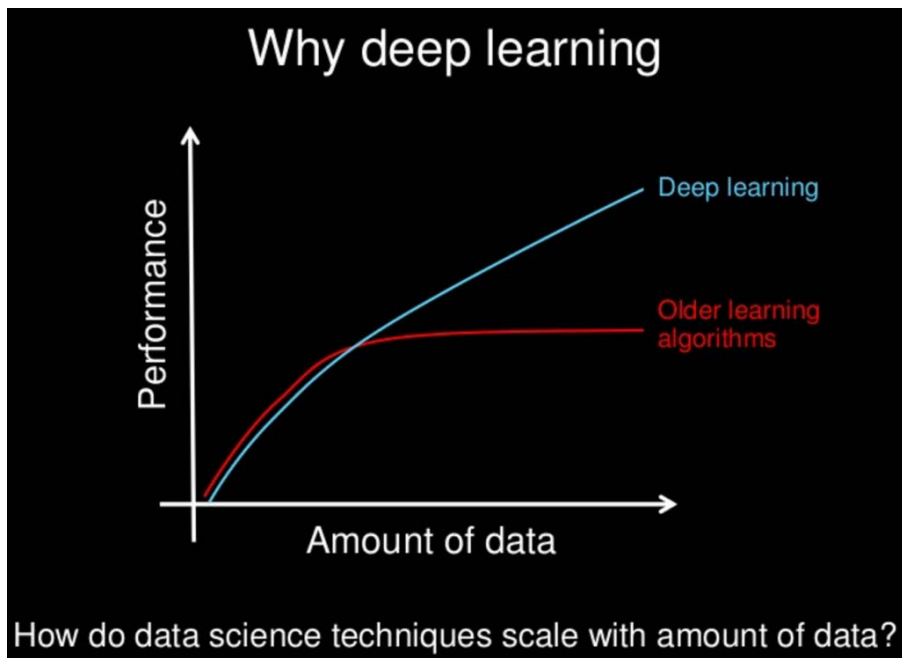
- Forward Propagation
- Backward Propagation

## 4. Deep Neural Networks

- Hidden Layers
- **Why Deep Neural Networks?**
- Types of DNNs
  - Multilayer Perceptron (ANN)
  - Convolutional Neural Network
  - Recurrent Neural Network
  - Generative Adversarial Networks (GANs)

## 5. Applications of Deep Neural Networks

# Why *Deep* Neural Networks?




- Data - Today, more than **2 trillion terabytes of data** is generated **each year**
- Hardware - We now have devices and **hardware capable of dealing with such volume of data**, eg. GPUs, TPUs, NPUs, etc.
- Performance - Deep Learning > Machine Learning almost universally now




# Why *Deep* Neural Networks - Performance

## Machine Learning



## Deep Learning



- Selecting features - **human limitation**
- **Deep learning** eliminates this requirement by doing **feature selection within the network internally**
- Often **still need to do manual preprocessing** to make DNN work

# Agenda

## 1. Neural Network Building Blocks

- What is a Neuron?
- Working of a Neuron
- Analogy with Human Brain
- Perceptron

## 2. Why we need Neural Networks

## 3. Working of a Neural Network

- Forward Propagation
- Backward Propagation

## 4. Deep Neural Networks

- Hidden Layers
- Why Deep Neural Networks?
- Types of DNNs
  - Multilayer Perceptron (ANN)
  - Convolutional Neural Network
  - Recurrent Neural Network
  - Generative Adversarial Networks (GANs)


## 5. Applications of Deep Neural Networks

# Types of *Deep* Neural Networks - Breakthroughs


## Deep Learning Categories

Main research areas and breakthroughs of DL


General Deep Learning  
Fully-Connected (FC)




2D/3D Image model  
CNN, FCN, etc.



1D Sequence Model  
RNN, LSTM, etc.



Others: unsupervised DL, reinforce Learning





# Agenda

## 1. Neural Network Building Blocks

- What is a Neuron?
- Working of a Neuron
- Analogy with Human Brain
- Perceptron

## 2. Why we need Neural Networks

## 3. Working of a Neural Network


- Forward Propagation
- Backward Propagation

## 4. Deep Neural Networks

- Hidden Layers
- Why Deep Neural Networks?
- Types of DNNs
  - **Multilayer Perceptron (ANN)**
  - Convolutional Neural Network
  - Recurrent Neural Network
  - Generative Adversarial Networks (GANs)

## 5. Applications of Deep Neural Networks

# Types of Neural Networks - Multilayer Perceptron



- Also known as **conventional Artificial Neural Network (ANN)** or **Vanilla Neural Network**
- Hyperparameters -
  - No. of hidden layers
  - No. of nodes in a layer
  - Learning rate for gradient descent, etc
- **Optimal combination** of hyperparameters - **experimentation**

# Agenda

## 1. Neural Network Building Blocks

- What is a Neuron?
- Working of a Neuron
- Analogy with Human Brain
- Perceptron

## 2. Why we need Neural Networks

## 3. Working of a Neural Network


- Forward Propagation
- Backward Propagation

## 4. Deep Neural Networks

- Hidden Layers
- Why Deep Neural Networks?
- Types of DNNs
  - Multilayer Perceptron (ANN)
  - **Convolutional Neural Network**
  - Recurrent Neural Network
  - Generative Adversarial Networks (GANs)

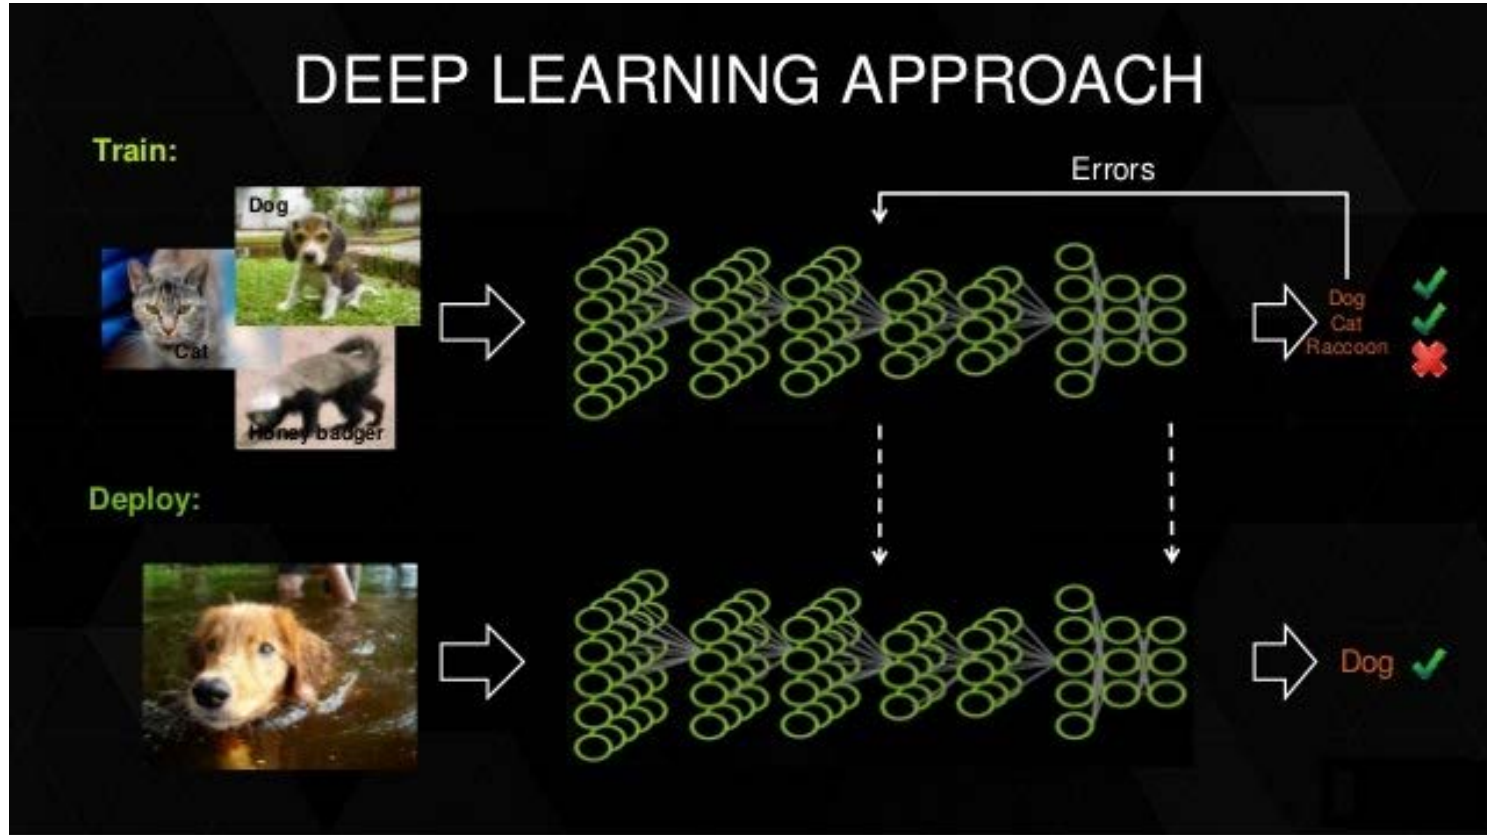
## 5. Applications of Deep Neural Networks

# Types of NN - Convolutional Neural Network (CNN)



- CNNs (by Geoffrey Hinton) are a type of NN **used often for image/video data**
- Use a **window to scan** over an image to **look for prominent features** in images
- Introduced **new layers** for better data processing like **pooling, dropout, etc.**

# Deep Neural Network - CNN as Eyes



# Agenda

## 1. Neural Network Building Blocks

- What is a Neuron?
- Working of a Neuron
- Analogy with Human Brain
- Perceptron

## 2. Why we need Neural Networks

## 3. Working of a Neural Network

- Forward Propagation
- Backward Propagation

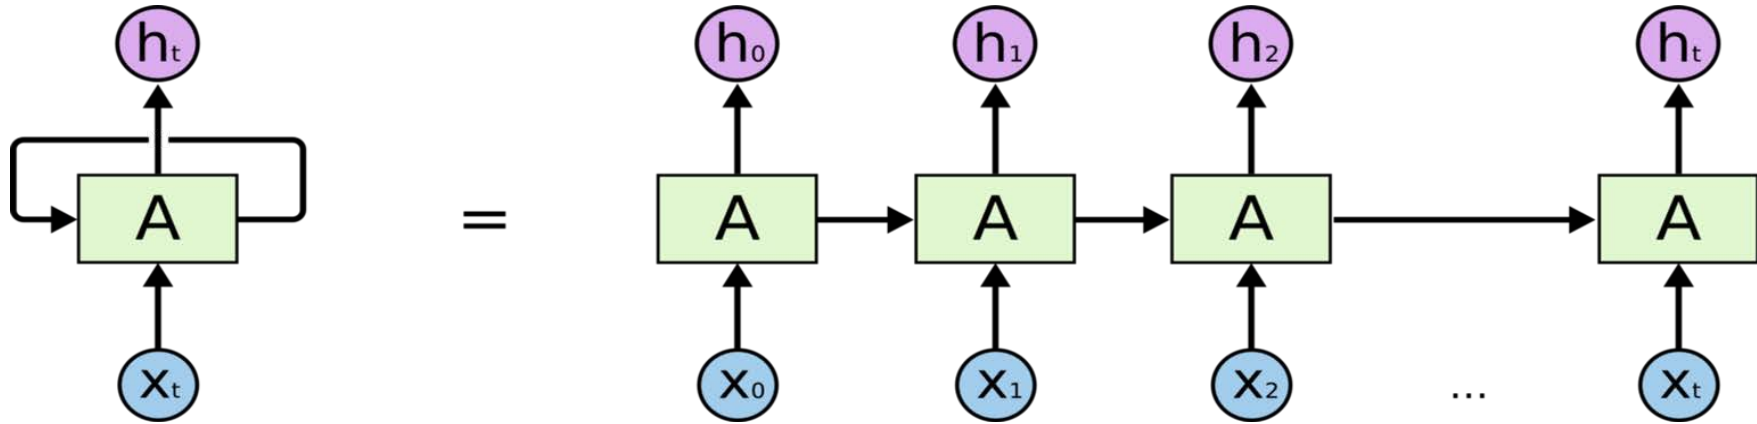
## 4. Deep Neural Networks

- Hidden Layers
- Why Deep Neural Networks?
- Types of DNNs
  - Multilayer Perceptron (ANN)
  - Convolutional Neural Network
  - **Recurrent Neural Network**
  - Generative Adversarial Networks (GANs)

## 5. Applications of Deep Neural Networks

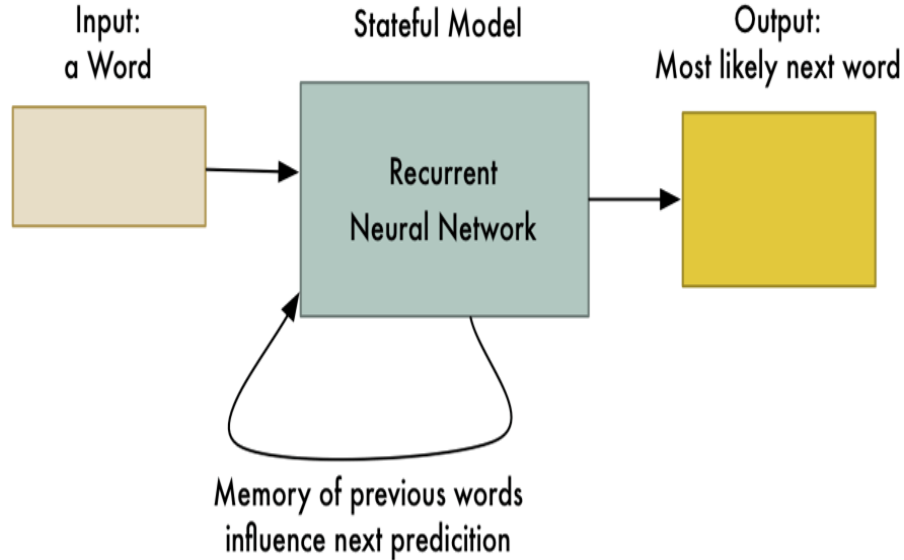


# Types of NN - Recurrent Neural Network (RNN)



- Use **internal memory to predict next state** making it ideal for audio, text and other **sequential data**
- Each layer has two inputs, the present and the recent past i.e. the **current layer output and the previous layer output**
- Remember the previous input states to give output for further states


# Deep Neural Network - RNN as Ears and Mouth



- RNNs are used to read and predict text
- They can also be used for identifying sounds such as song lyrics
- We also use them to generate text, as well as speech



# RNN - Working Example - Predict next word in sentence



- Predict the next word based on previous
- Output at each time-step - list of next probable words
- Pick most probable

# Agenda

## 1. Neural Network Building Blocks

- What is a Neuron?
- Working of a Neuron
- Analogy with Human Brain
- Perceptron

## 2. Why we need Neural Networks

## 3. Working of a Neural Network

- Forward Propagation
- Backward Propagation


## 4. Deep Neural Networks

- Hidden Layers
- Why Deep Neural Networks?
- Types of DNNs
  - Multilayer Perceptron (ANN)
  - Convolutional Neural Network
  - Recurrent Neural Network
  - **Generative Adversarial Networks (GANs)**

## 5. Applications of Deep Neural Networks

# Types of NN - Generative Adversarial Networks (GANs)

## Generative Adversarial Networks (GANs)

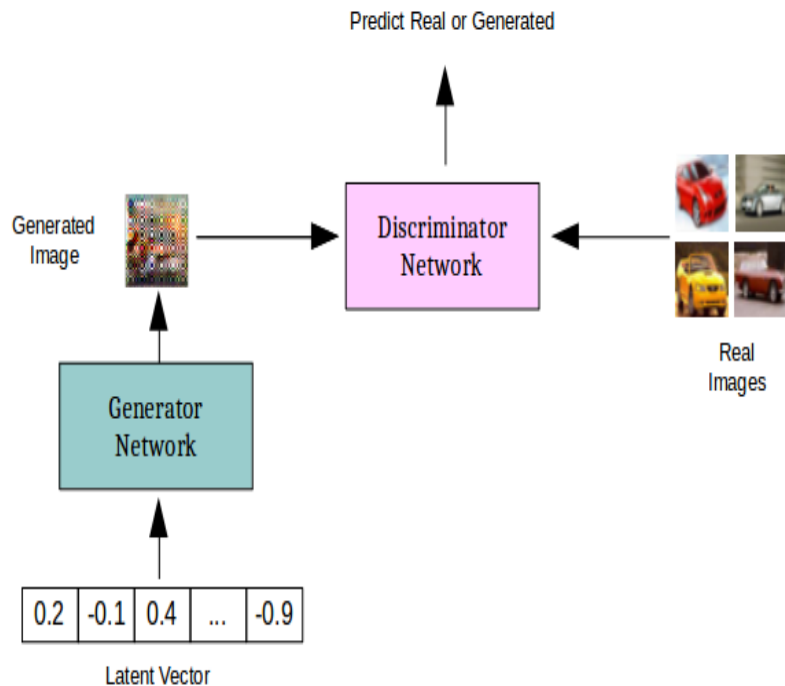


Try to develop fake currency to fool police

Use methods to identify the fake currency

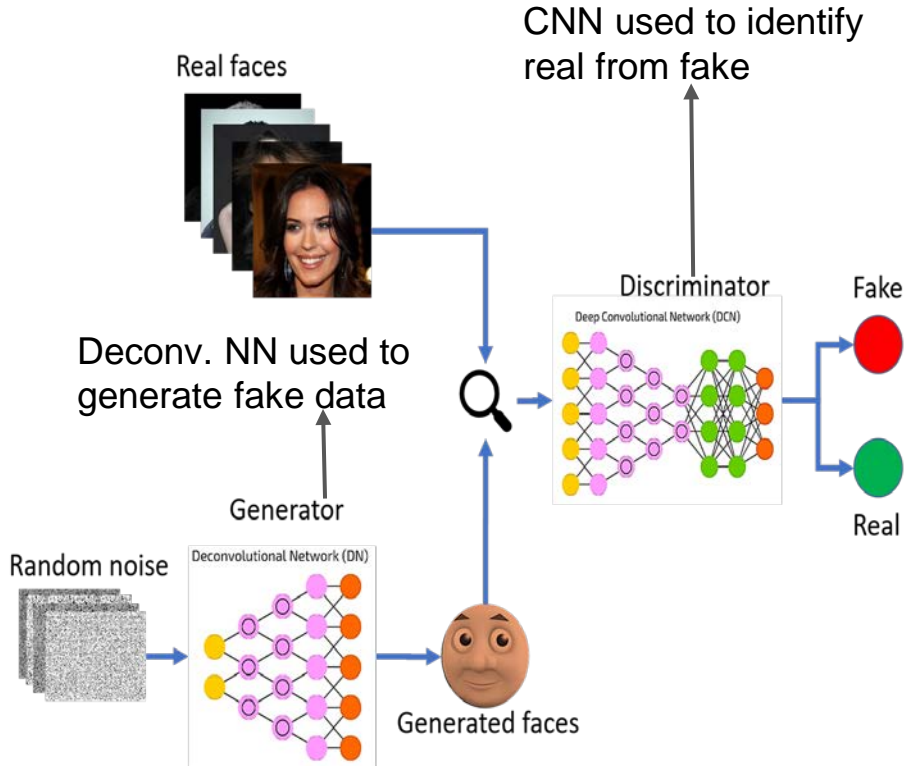
- Deal with data generation and imitation
- Two networks -
  - **Generator** - generates new data
  - **Discriminator** - evaluates authenticity
- Generator : Counterfeiter as Discriminator : Police

# Generative Adversarial Networks (GANs)



- **Generator** - takes random data vector (noise), transforms it to target data (eg. cars)
- **Discriminator** - Tries to differentiate the generated data from real, outputs a probability

# Generative Adversarial Networks (GANs) - Working



- If discriminator **fooled** - **discriminator improved** to identify better
- If discriminator **catches forgery** - **generator makes more authentic data**

# Agenda

## 1. Neural Network Building Blocks

- What is a Neuron?
- Working of a Neuron
- Analogy with Human Brain
- Perceptron

## 2. Why we need Neural Networks

## 3. Working of a Neural Network

- Forward Propagation
- Backward Propagation


## 4. Deep Neural Networks

- Hidden Layers
- Why Deep Neural Networks?
- Types of DNNs
  - Multilayer Perceptron (ANN)
  - Convolutional Neural Network
  - Recurrent Neural Network
  - Generative Adversarial Networks (GANs)


## 5. Applications of Deep Neural Networks

# Applications of Deep Neural Networks

## 1. Image Colourization



## 2. Machine Translation



## 3. Auto Handwriting Generation


Machine Learning Mastery

Machine Learning Mastery

Machine Learning Mastery



## 4. Object Classification and Detection



## 5. Image Captioning



"man in black shirt is playing guitar."



"construction worker in orange safety vest is working on road."



"two young girls are playing with lego toy."



"girl in pink dress is jumping in air."



"black and white dog jumps over bar."



"young girl in pink shirt is swinging on swing."

## 6. Self-Driving Cars

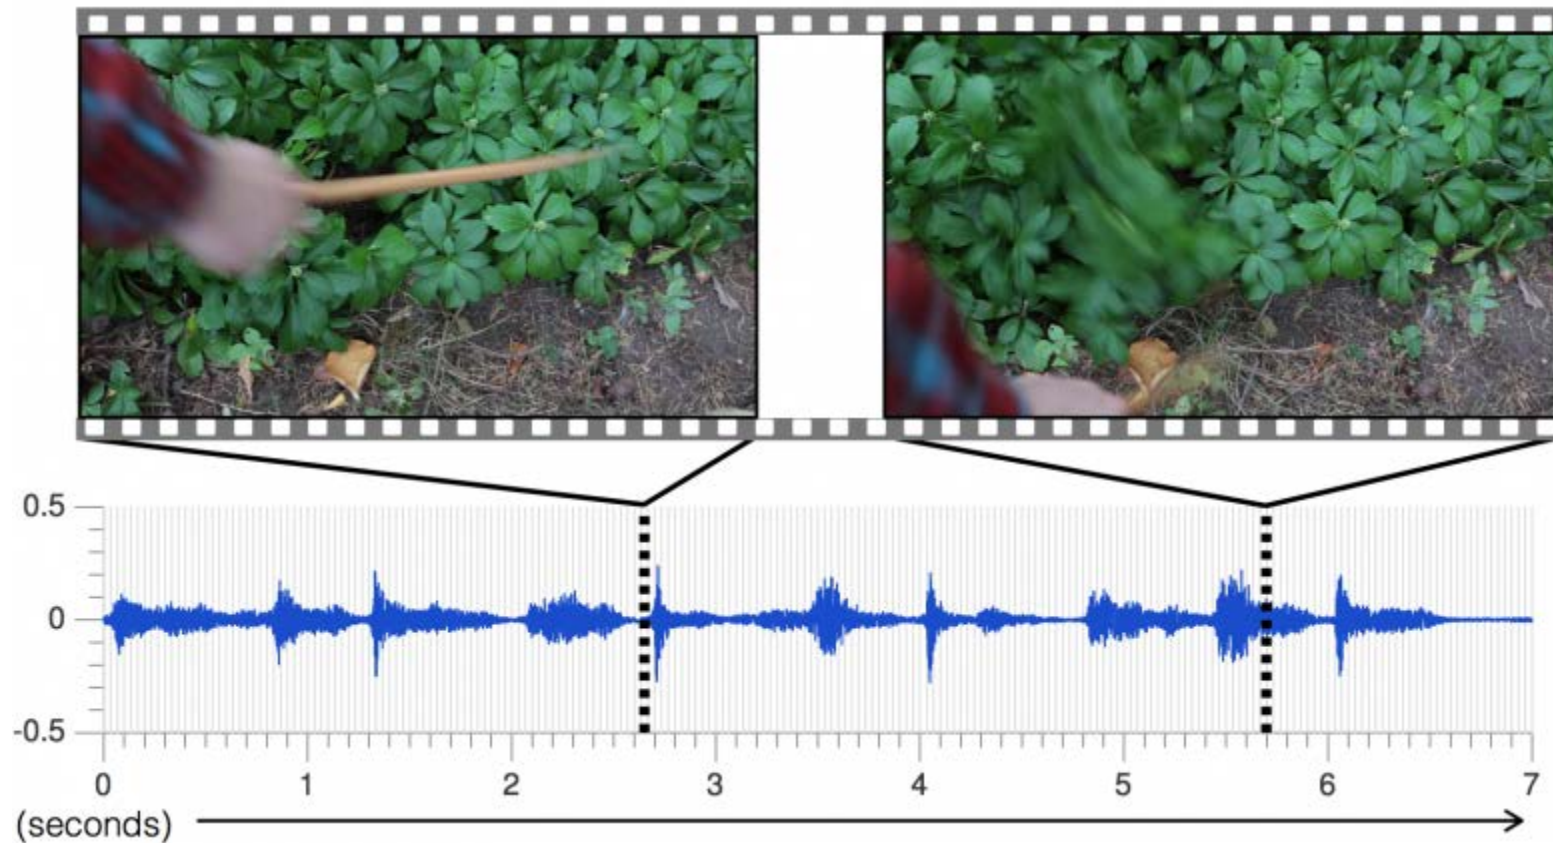


## 7. Style Transfer





## 8. Adding Sounds to Silent Movies



## 9. Image Handling

JPEG Artifacts removal



Corrupted



Deep image prior

Inpainting



Corrupted



Deep image prior

Inpainting



Corrupted



Deep image prior

Super-resolution



Corrupted



Deep image prior

Denoising



Corrupted



Deep image prior

Inpainting



Corrupted



Deep image prior

Deep Image  
Prior - a GAN

## 10. Sentiment Analysis

# SENTIMENT ANALYSIS

---



### NEGATIVE

Totally dissatisfied with the service. Worst customer care ever.



### NEUTRAL

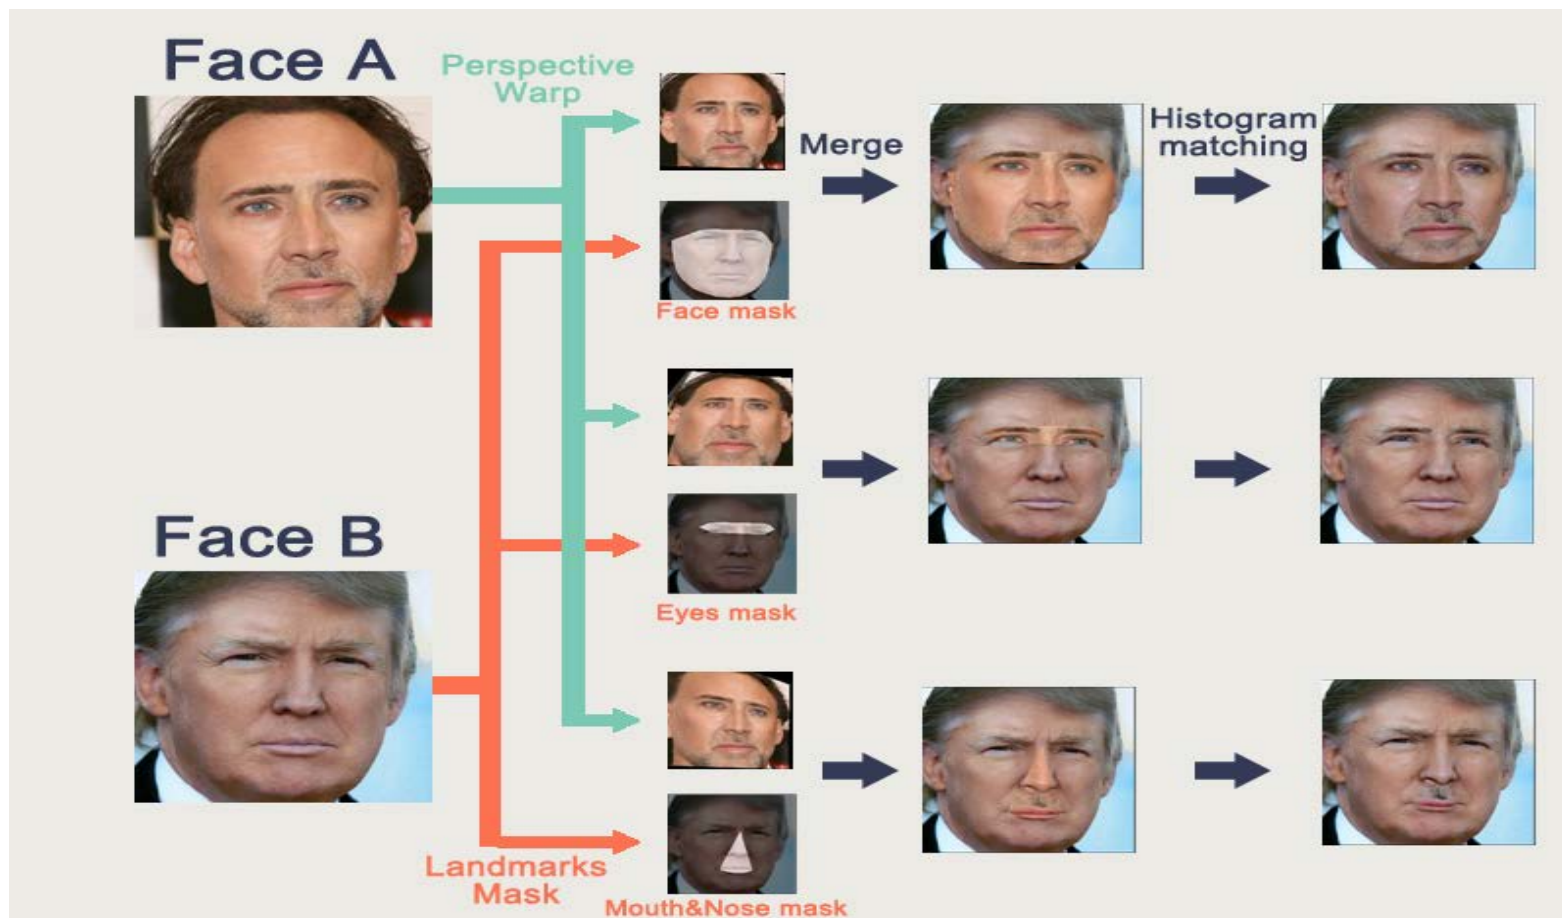
Good Job but I will expect a lot more in future.



### POSITIVE

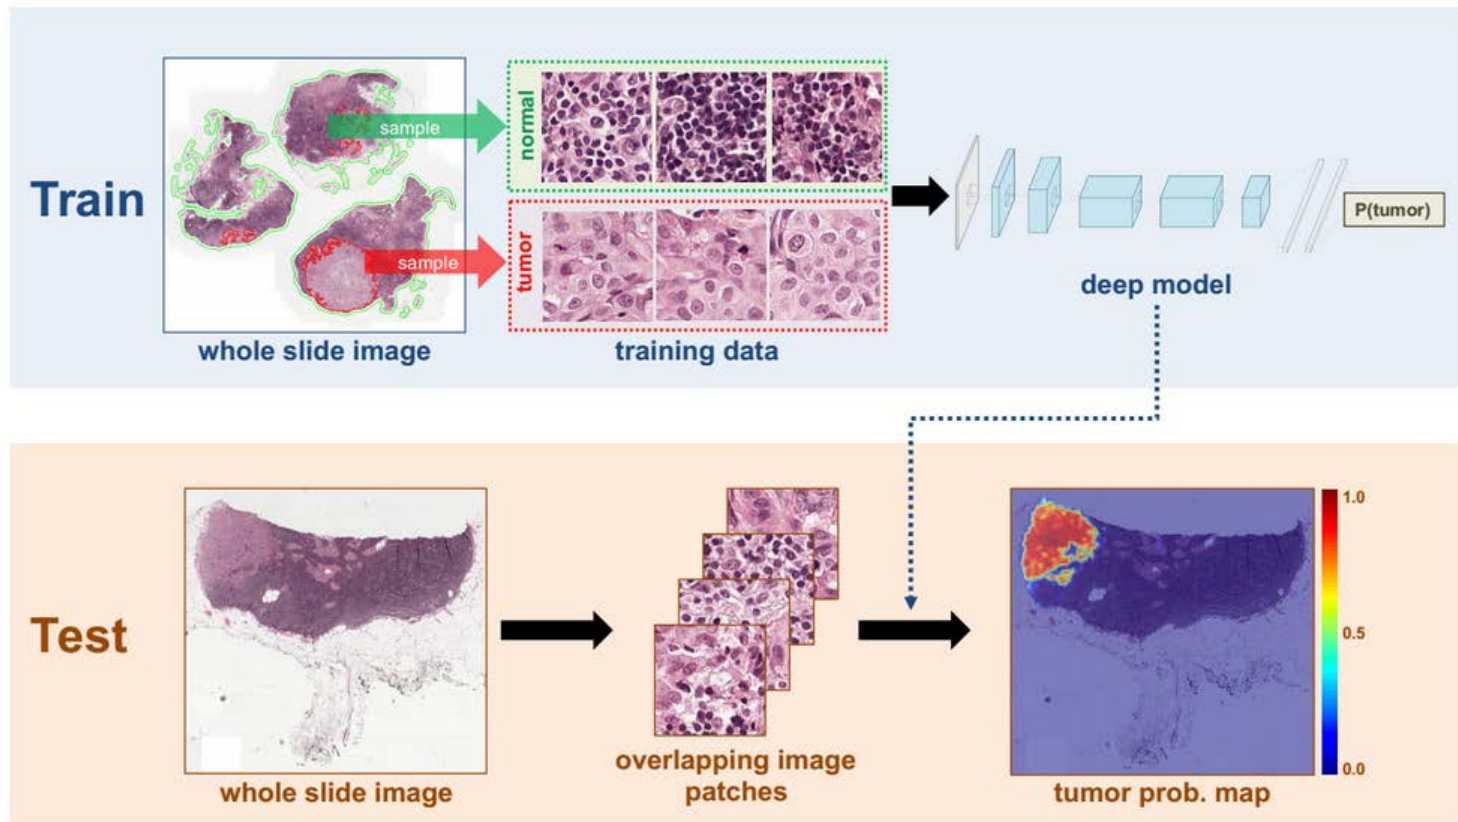
Brilliant effort guys! Loved Your Work.

## 11. Face Swapping





## 12. Cancer Detection and Healthcare



And many more...